

Lecture 5 :

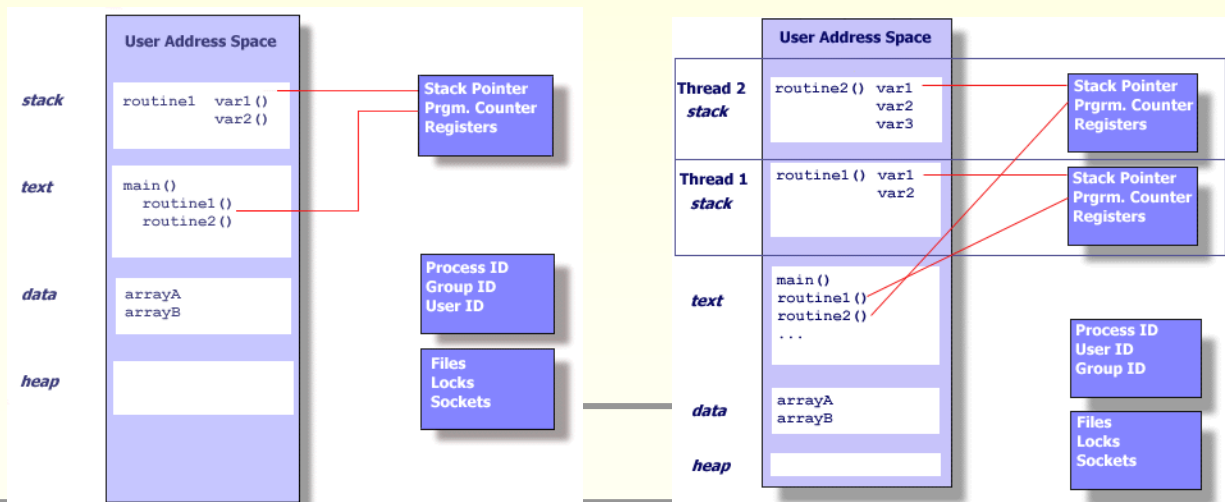
Pthread Programming

What is Thread?

- Independent stream of instructions executed simultaneously by OS
- Multithreaded program
 - A main program contains a number of procedures that are scheduled to run simultaneously and independently by OS

Processes and Threads

- Threads share resources of a process
 - Changes made by one thread affect other threads
 - Two pointers having the same value point to the same data
 - Reading and writing to the same memory location is possible, and therefore require explicit synchronization by programmer
- Processes don't share resources



Thread Properties

- Exists within a process and uses the process resources
- Has its own independent flow of control as long as its parent process exists and the OS supports it
- Duplicates only the essential resources it needs to be independently schedulable
- May share the process resources with other threads that act equally independently (and dependently)
- Dies if the parent process dies – or something similar
- Is "lightweight" because most of the overhead has already been accomplished through the creation of its process.
- All threads within a process share same address space.
- Therefore, inter-thread communication is more efficient than inter-process communication

pthread

■ pthread

- POSIX thread
- Standardized C language threads for UNIX
- For Portability
- Working in **shared memory multiprocessor**

■ Why pthreads?

- Performance gains
- Requires fewer system resources than process
 - Compare fork() and pthread_create() : 10~50 times

fork() vs pthread_create()

- Performance (50,000 fork or pthread_create)

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

Pthreads API

- Three groups
 - Thread Management
 - Thread creation, and destruction
 - Mutexes (mutual exclusion)
 - synchronization
 - Conditional Variables
 - Communication between threads that share a mutex

Thread Management

- `pthread_create (thread,attr,start_routine,arg)`
- `pthread_exit (status)`
- `pthread_cancel (thread)`

Thread Creation

```
int pthread_create(  
    pthread_t *restrict thread,  
    const pthread_attr_t *restrict attr,  
    void *(*start_routine)(void *),  
    void *restrict arg);
```

- Creates a new thread and makes it executable.
- The creating process (or thread) must provide a location for storage of the thread id.
- The third parameter is just the name of the function for the thread to run.
- The last parameter is a pointer to the arguments.

Thread Creation

- When a new thread is created, it runs concurrently with the creating process.
- When creating a thread, you indicate which function the thread should execute.
- Thread handle returned via **pthread_t** structure
- Specify **NULL** to use default attributes
- Single argument sent to the function
- If no arguments to function, specify **NULL**
- Check error codes!

Thread Termination

- There are several ways in which a pthread may be terminated:
 - The thread returns from its starting routine (the main routine for the initial thread).
 - The thread makes a call to the **pthread_exit** subroutine.
 - The thread is canceled by another thread via the **pthread_cancel** routine
 - The entire process is terminated due to making a call to either the `exec()` or `exit()`
 - If `main()` finishes first, without calling **pthread_exit** explicitly itself

Thread Termination

```
void pthread_exit(void *value_ptr);
```

- Typically, the pthread_exit() routine is called to quit the thread.
- If main() finishes before the threads it has created, and exits with pthread_exit(), the other threads will continue to execute. Otherwise, they will be automatically terminated when main() finishes.
- The programmer may optionally specify a termination status, which is stored as a void pointer for any thread that may join the calling thread.
- Cleanup: the pthread_exit() routine does not close files; any files opened inside the thread will remain open after the thread is terminated.

Thread Cancellation

- One thread can request that another exit with `pthread_cancel`
- **`int pthread_cancel(pthread_t thread) ;`**
- The `pthread_cancel` returns after making the request.

Example 1: Thread Creation

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid) {
    int tid;
    tid = (int)threadid;
    printf("Hello World! It's me, thread #%d!\n", tid);
    pthread_exit(NULL);
}

int main (int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    int rc, t;
    for(t=0; t<NUM_THREADS; t++){
        printf("In main: creating thread %d\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc) {
            printf("ERROR code is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Example 2: Passing Parameters to Thread

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

void *PrintHello(void *ptr)
{
    char *filename;
    int j;
    filename = (char *) ptr;
    while (1) {
        printf("Hello World! It's me, thread %s!\n", filename);
        sleep(1);
    }
    pthread_exit(NULL);
}

int main (int argc, char *argv[])
{
    pthread_t thread[100];
    int err_code, i=0;
    char *filename;

    printf ("Enter thread name at any time to create thread\n");

    while (1) {
        filename = (char *) malloc (80*sizeof(char));
        scanf ("%s", filename);
        printf("In main: creating thread %d\n", i);
        err_code = pthread_create(&thread[i], NULL, PrintHello, (void *)filename);
        if (err_code){
            printf("ERROR code is %d\n", err_code);
            exit(-1);
        } else i++;
    }
    pthread_exit(NULL);
}
```

Example 3: Files

```
void *PrintHello(void *ptr)
{
    FILE *file;
    char *filename;
    char textline[100];
    int j;

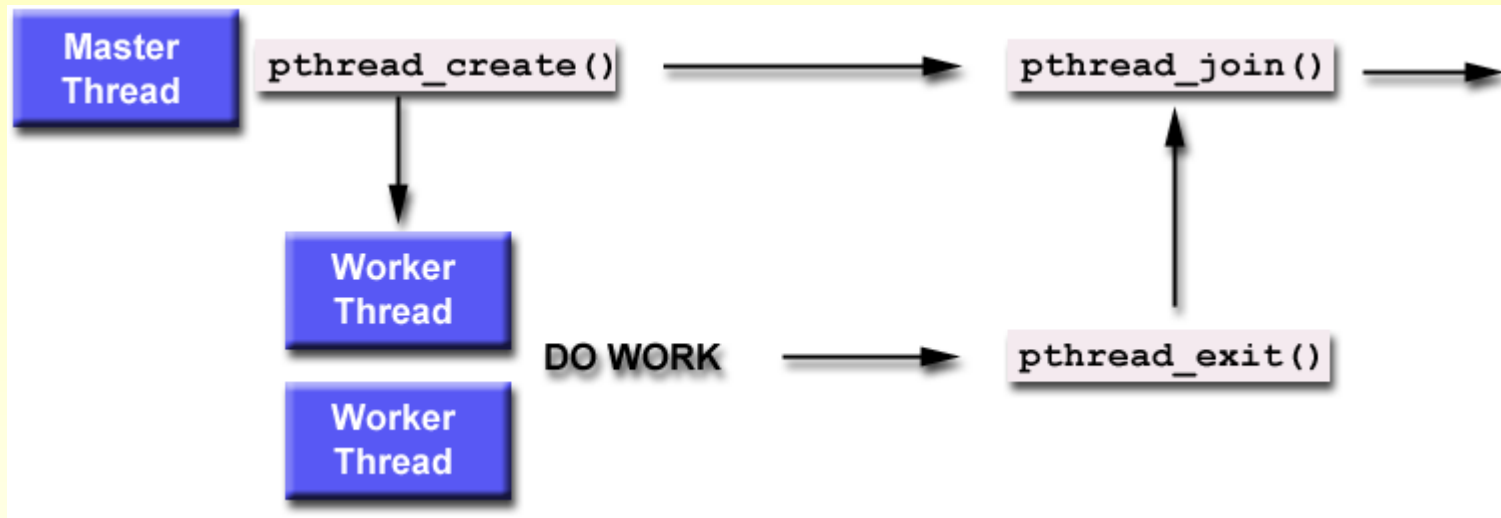
    filename = (char *) ptr;
    printf("Hello World! Opening %s!\n", filename);
    file = fopen(filename, "r");
    if (file != NULL) {
        while(fscanf(file, "%s\n", textline) != EOF) printf ("%s\n", textline);
    }
    fclose(file);
}
```


JOINING

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- The pthread_join() subroutine blocks the calling thread until the specified thread terminates.
- The programmer is able to obtain the target thread's termination return status if it was specified in the target thread's call to pthread_exit().
- A joining thread can match one pthread_join() call. It is a logical error to attempt multiple joins on the same thread.

JOINING



Example 4: JOIN

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <math.h>
#define NUM_THREADS 4

void *BusyWork(void *t)
{
    int i;
    long tid;
    double result=0.0;
    tid = (long)t;
    printf("Thread %ld starting...\n",tid);
    for (i=0; i<1000000; i++)
    {
        result = result + sin(i) * tan(i);
    }
    printf("Thread %ld done. Result = %e\n",tid, result);
    pthread_exit((void*) t);
}
```

```
int main (int argc, char *argv[])
{
    pthread_t thread[NUM_THREADS];
    int rc;
    long t;
    void *status;

    for(t=0; t<NUM_THREADS; t++) {
        printf("Main: creating thread %ld\n", t);
        rc = pthread_create(&thread[t], &attr, BusyWork, (void *)t);
        if (rc) {
            printf("ERROR; return code from pthread_create()
                    is %d\n", rc);
            exit(-1);
        }
    }

    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_join(thread[t], &status);
        if (rc) {
            printf("ERROR; return code from pthread_join() is %d\n", rc);
            exit(-1);
        }
        printf("Main: completed join with thread %ld having a status
                of %ld\n", t, (long)status);
    }
    printf("Main: program completed. Exiting.\n");
    pthread_exit(NULL);
}
```

Mutexes

- Mutual Exclusion
- implementing thread synchronization and protecting shared data when multiple writes occur.
- A mutex variable acts like a "lock" protecting access to a shared data resource
 - only one thread can lock (or own) a mutex variable at any given time
- Used for preventing race condition
- When several threads compete for a mutex, the losers block at that call – an unblocking call is available with "trylock" instead of the "lock" call.

Mutex Routines

- `pthread_mutex_init (mutex, attr)`
- `pthread_mutex_destroy (mutex)`
- Mutex variables must be declared with type **`pthread_mutex_t`**, and must be initialized before they can be used

Locking/Unlocking Mutexes

- **pthread_mutex_lock (mutex)**
 - acquire a lock on the specified *mutex* variable
- **pthread_mutex_trylock (mutex)**
 - attempt to lock a mutex. However, if the mutex is already locked, the routine will return immediately with a "busy" error code
- **pthread_mutex_unlock (mutex)**
 - unlock a mutex if called by the owning thread

User's Responsibility for Using Mutex

- When protecting shared data, it is the programmer's responsibility to make sure every thread that needs to use a mutex does so.
- For example, if 3 threads are updating the same data, but only one or two use a mutex, the data can still be corrupted.

Thread 1	Thread 2	Thread 3
Lock	Lock	
$A = 2$	$A = A + 1$	$A = A * B$
Unlock	Unlock	

Example 5: Mutexes

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

typedef struct
{
    double      *a;
    double      *b;
    double      sum;
    int         veclen;
} DOTDATA;

#define NUMTHRDS 4
#define VECLEN 100000
    DOTDATA dotstr;
    pthread_t callThd[NUMTHRDS];
    pthread_mutex_t mutexsum;

void *dotprod(void *arg)
{
    int i, start, end, len ;
    long offset;
    double mysum, *x, *y;
    offset = (long)arg;

    len = dotstr.vecLen;
    start = offset*len;
    end   = start + len;
    x = dotstr.a;    y = dotstr.b;
    mysum = 0;
    for (i=start; i<end ; i++) mysum += (x[i] * y[i]);
    pthread_mutex_lock (&mutexsum);
    dotstr.sum += mysum;
    printf("Thread %ld did %d to %d: mysum=%f global sum=%f\n",offset,start,end,mysum,dotstr.sum);
    pthread_mutex_unlock (&mutexsum);
    pthread_exit((void*) 0);
}
```

Example 5: Mutexes

```
int main (int argc, char *argv[])
{
    long i;
    double *a, *b;
    void *status;

    a = (double*) malloc (NUMTHRDS*VECLEN*sizeof(double));
    b = (double*) malloc (NUMTHRDS*VECLEN*sizeof(double));

    for (i=0; i<VECLEN*NUMTHRDS; i++) {
        a[i]=1;
        b[i]=a[i];
    }

    dotstr.vecLEN = VECLen;
    dotstr.a = a;
    dotstr.b = b;
    dotstr.sum=0;

    pthread_mutex_init(&mutexsum, NULL);

    for(i=0;i<NUMTHRDS;i++) pthread_create(&callThd[i], &attr, dotprod, (void *)i);

    for(i=0;i<NUMTHRDS;i++) pthread_join(callThd[i], &status);
    printf ("Sum =  %f \n", dotstr.sum);

    free (a);
    free (b);
    pthread_mutex_destroy(&mutexsum);
    pthread_exit(NULL);
}
```

Condition Variables

- another way for threads to synchronize
- mutexes
 - synchronization by controlling thread access to data
- condition variables
 - synchronization based upon the actual value of data.
 - Without condition variables, the programmer would need to have threads continually polling (possibly in a critical section), to check if the condition is met
 - always used in conjunction with a mutex lock

Condition Variables Routines

- `pthread_cond_init (condition,attr)`
- `pthread_cond_destroy (condition)`
- Condition variables must be declared with type `pthread_cond_t`, and must be initialized before they can be used.
- *attr* is used to set condition variable attributes.(NULL: defaults)
- `pthread_cond_destroy()` should be used to free a condition variable that is no longer needed.

Condition Variables Routines

- **pthread_cond_wait (condition,mutex)**
 - blocks the calling thread until the specified *condition* is signalled.
 - This routine should be called while *mutex* is locked
 - will automatically release the mutex lock while it waits
 - After signal is received and thread is awakened, *mutex* will be automatically locked for use
- **pthread_cond_signal (condition)**
 - signal (or wake up) another thread which is waiting on the condition variable.
 - It is a logical error to call pthread_cond_signal() before calling pthread_cond_wait().
- **pthread_cond_broadcast (condition)**
 - should be used instead of pthread_cond_signal() if more than one thread is in a blocking wait state.

Example 6: Control Variables

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

#define NUM_THREADS 3
#define TCOUNT 10
#define COUNT_LIMIT 12

int count = 0;
int thread_ids[3] = {0,1,2};
pthread_mutex_t count_mutex;
pthread_cond_t count_threshold_cv;

void *inc_count(void *t)
{
    int i;
    long my_id = (long)t;

    for (i=0; i<TCOUNT; i++) {
        pthread_mutex_lock(&count_mutex);
        count++;

        if (count == COUNT_LIMIT) {
            pthread_cond_signal(&count_threshold_cv);
            printf("inc_count(): thread %ld, count = %d Threshold reached.\n", my_id, count);
        }
        printf("inc_count(): thread %ld, count = %d, unlocking mutex\n", my_id, count);
        pthread_mutex_unlock(&count_mutex);

        sleep(1);
    }
    pthread_exit(NULL);
}
```

```

void *watch_count(void *t)
{
    long my_id = (long)t;

    printf("Starting watch_count(): thread %ld\n", my_id);
    pthread_mutex_lock(&count_mutex);
    while (count < COUNT_LIMIT) {
        pthread_cond_wait(&count_threshold_cv, &count_mutex);
        printf("watch_count(): thread %ld Condition signal received.\n", my_id);
        count += 125;
        printf("watch_count(): thread %ld count now = %d.\n", my_id, count);
    }
    pthread_mutex_unlock(&count_mutex);
    pthread_exit(NULL);
}

```

```

int main (int argc, char *argv[])
{
    int i, rc;
    long t1=1, t2=2, t3=3;
    pthread_t threads[3];
    pthread_mutex_init(&count_mutex, NULL);
    pthread_cond_init (&count_threshold_cv, NULL);
    pthread_create(&threads[0], NULL, watch_count, (void *)t1);
    pthread_create(&threads[1], NULL, inc_count, (void *)t2);
    pthread_create(&threads[2], NULL, inc_count, (void *)t3);

    for (i=0; i<NUM_THREADS; i++) pthread_join(threads[i], NULL);
    printf ("Main(): Waited on %d threads. Done.\n", NUM_THREADS);
    pthread_mutex_destroy(&count_mutex);
    pthread_cond_destroy(&count_threshold_cv);
    pthread_exit(NULL);
}

```